

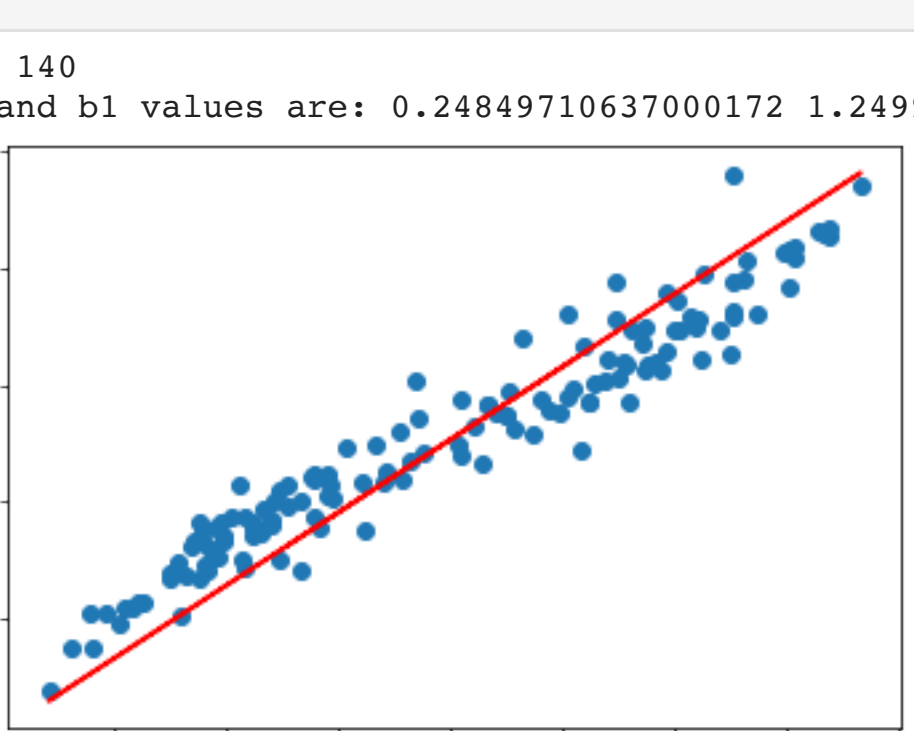
```
In [1]: import numpy as np
import sympy as sym
import matplotlib.pyplot as plt
import matplotlib.pyplot as plt
npfile = np.load('data-2.npy')
x = npfile[:,0].reshape(140,1)
y = npfile[:,1].reshape(140,1)
print(len(x),len(y))
iterations = 5000
plt.scatter(x,y)
b0 = 0
b1 = 0
d_b0 = 0
d_b1 = 0
X = npfile[:,0].reshape(140,1)
# X = np.concatenate((np.ones((X.shape[0], 1)), X), axis=1)
m = len(y)
alpha = 0.0001
def cost_function(b0,b1,X,y,iterations=5000):
    J = (1/(2*m))*sum((-b0-(b1*X)+y)**2)
    return J

def gradient(b0,b1,X,y):
    d_b0 = -(1/m)*sum(y-b0-(b1*X))
    d_b1 = -(1/m)*sum(X*(y-b0-(b1*X)))
    return d_b0, d_b1


def gradient_descent(b0,b1,X,y,alpha,iterations=5000):
    cost = []
    for i in range(iterations):
        [d_b0, d_b1] = gradient(b0,b1,X,y)
        cost.append(cost_function(b0,b1,X,y)[0])
        b0 = b0-(alpha*d_b0)
        b1 = b1-(alpha*d_b1)
    return b0[0],b1[0],cost
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha)
print('b0 and b1 values are:',b0,b1)

Y = b0+b1*X
plt.plot(X,Y,'r')
plt.show()
# print(cost)
f = np.linspace(0,iterations,5000)
plt.plot(f[:100],cost[:100])

140 140
b0 and b1 values are: 0.24849710637000172 1.249979368235588
```



Out[1]: [




```
2(b)

In [2]: alpha = 0.0002
b0 = 0
b1 = 0
iterations = 5000
f = np.linspace(0,iterations,5000)
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha)
print('b0 and b1 values are:',b0,b1)

plt.plot(f[:100],cost[:100])

b0 and b1 values are: 0.4736728136091704 1.246310709297146
```

Out[2]: [




```
In [3]: alpha = 0.0006
b0 = 0
b1 = 0
iterations = 5000
f = np.linspace(0,iterations,10)
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha)
print('b0 and b1 values are:',b0,b1)

plt.plot(f,cost[:10])

<ipython-input-1-ea25a796d343>:19: RuntimeWarning: overflow encountered in add
J = (1/(2*m))*sum((-b0-(b1*X)+y)**2)
<ipython-input-1-ea25a796d343>:19: RuntimeWarning: overflow encountered in square
J = (1/(2*m))*sum((-b0-(b1*X)+y)**2)
<ipython-input-1-ea25a796d343>:24: RuntimeWarning: overflow encountered in add
d_b1 = -(1/m)*sum(X*(y-b0-(b1*X)))
<ipython-input-1-ea25a796d343>:34: RuntimeWarning: invalid value encountered in subtract
b1 = b1-(alpha*d_b1)
b0 and b1 values are: nan nan
```

Out[3]: [



alpha = 0.01 b0 = 0 b1 = 0 iterations = 5000 f = np.linspace(0,iterations,10) [b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha) print('b0 and b1 values are:',b0,b1)


plt.plot(f,cost[:10])

```
In [4]: alpha = 0.01
b0 = 0
b1 = 0
iterations = 5000
f = np.linspace(0,iterations,10)
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha)
print('b0 and b1 values are:',b0,b1)

plt.plot(f,cost[:10])

<ipython-input-1-ea25a796d343>:19: RuntimeWarning: overflow encountered in square
J = (1/(2*m))*sum((-b0-(b1*X)+y)**2)
<ipython-input-1-ea25a796d343>:19: RuntimeWarning: overflow encountered in add
J = (1/(2*m))*sum((-b0-(b1*X)+y)**2)
<ipython-input-1-ea25a796d343>:24: RuntimeWarning: overflow encountered in add
d_b1 = -(1/m)*sum(X*(y-b0-(b1*X)))
<ipython-input-1-ea25a796d343>:34: RuntimeWarning: invalid value encountered in subtract
b1 = b1-(alpha*d_b1)
b0 and b1 values are: nan nan
```

Out[4]: [




```
In [5]: alpha = 10
b0 = 0
b1 = 0
iterations = 5000
f = np.linspace(0,iterations,10)
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha)
print('b0 and b1 values are:',b0,b1)

plt.plot(f,cost[:10])

<ipython-input-1-ea25a796d343>:19: RuntimeWarning: overflow encountered in square
J = (1/(2*m))*sum((-b0-(b1*X)+y)**2)
<ipython-input-1-ea25a796d343>:23: RuntimeWarning: overflow encountered in add
d_b0 = -(1/m)*sum(y-b0-(b1*X))
<ipython-input-1-ea25a796d343>:24: RuntimeWarning: overflow encountered in multiply
d_b1 = -(1/m)*sum(X*(y-b0-(b1*X)))
<ipython-input-1-ea25a796d343>:33: RuntimeWarning: invalid value encountered in subtract
b0 = b0-(alpha*d_b0)
<ipython-input-1-ea25a796d343>:34: RuntimeWarning: invalid value encountered in subtract
b1 = b1-(alpha*d_b1)
b0 and b1 values are: nan nan
```

Out[5]: [

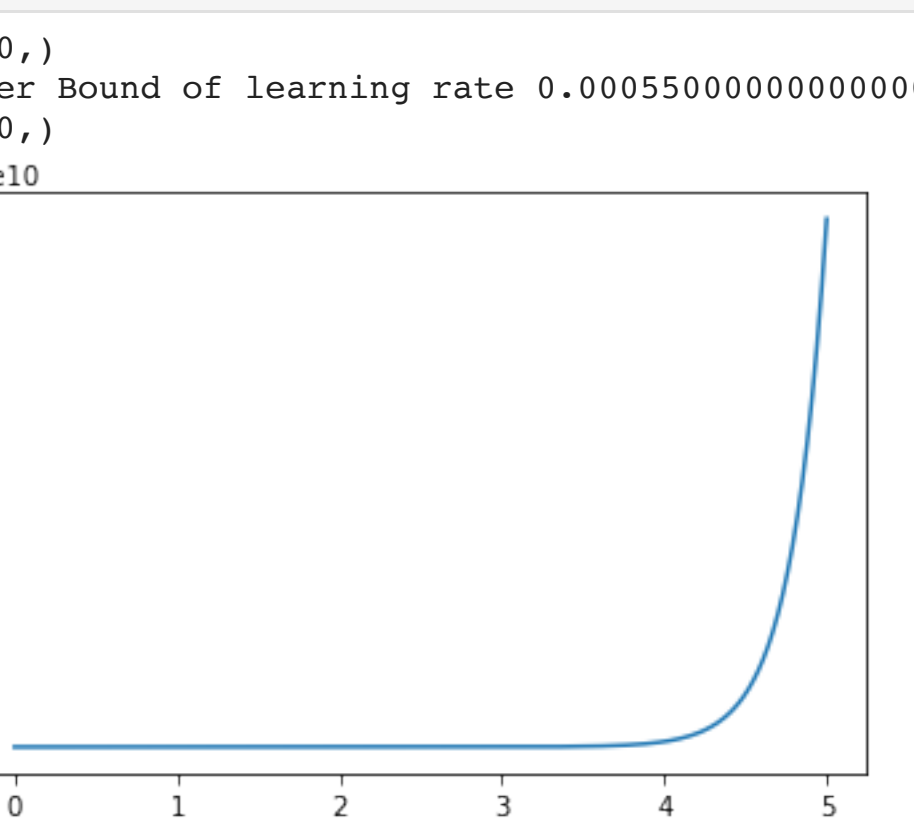


2(c)
The best learning rate is 0.0001 since it converges and converges quicker than the other learning rates


```
2(e)

In [6]: b0 = 0
b1 = 0
cost_l = []
iterations = 5000
f = np.linspace(0,5,500)
print(np.shape(f))
alpha = 0.000001
for i in range(iterations):
    alpha+=0.000001
    [b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha,iterations = 500)
    cost_l = cost
    if cost[-1]>=6000:
        break
alpha = alpha-0.000001
print('Upper Bound of learning rate',alpha)
print(np.shape(cost))
plt.plot(f,cost)
plt.show()
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha,iterations = 500)
plt.plot(f,cost)

(500,)
Upper Bound of learning rate 0.00055000000000000069
(500,)
```



Out[6]: [



2(f)

```
In [7]: npfile = np.load('data-2.npy')
x = npfile[:,0].reshape(140,1)
y = npfile[:,1].reshape(140,1)


a = np.arange(len(X))
np.random.shuffle(a)
np.random.seed(20)
print(len(x),len(y))
iterations = 5000
plt.scatter(x,y)
b0 = 0
b1 = 0
d_b0 = 0
d_b1 = 0
X = np.array(npfile[:,0].reshape(140,1))
X = np.concatenate((np.ones((X.shape[0], 1)), X), axis=1)
m = 20
alpha = 0.0005
def cost_function(b0,b1,X,y,iterations=5000):
    J = (1/(m))*sum((-b0-(b1*X)+y)**2)
    # print(J)
    return J

def gradient(b0,b1,X,y):
    d_b0 = -(1/m)*sum(y-b0-(b1*X))
    d_b1 = -(1/m)*sum(X*(y-b0-(b1*X)))
    # breakpoint()
    return d_b0, d_b1

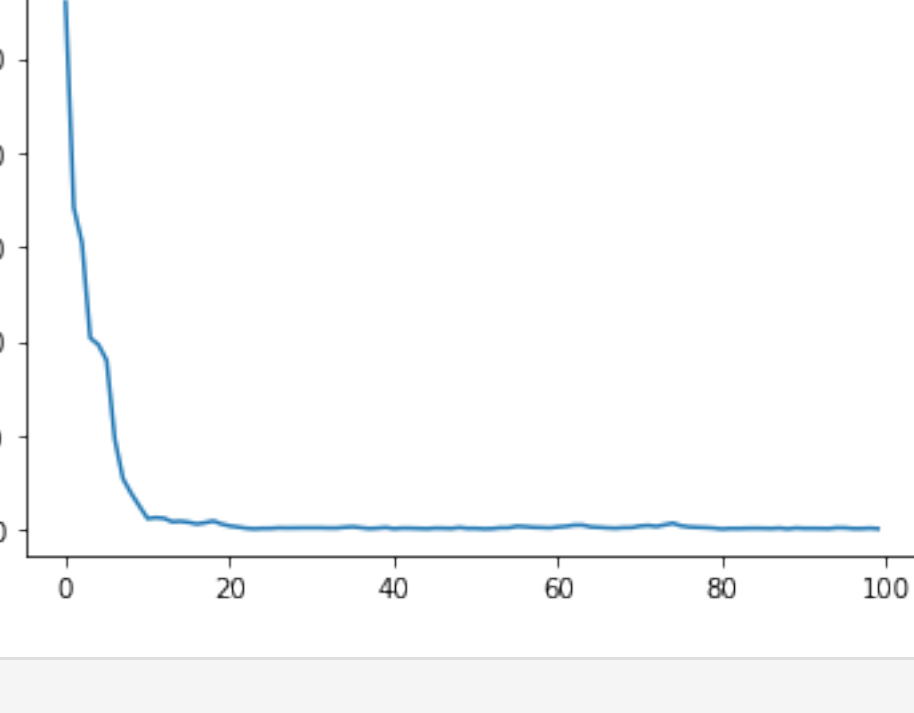
def gradient_descent(b0,b1,X,y,alpha,iterations=5000):
    cost = []
    for i in range(iterations):
        np.random.shuffle(a)
        for j in range(0,140,20):
            # print(a)
            indices = a[j:j+m]
            print(indices)
            # print(x_1)
            [d_b0, d_b1] = gradient(b0,b1,X[indices],y[indices])
            # print(np.shape(x_1))
            cost.append(cost_function(b0,b1,X[indices],y[indices]))
            # print(cost)
            b0 = b0-(alpha*d_b0)
            b1 = b1-(alpha*d_b1)
            # breakpoint()
    return b0,b1,cost
[b0, b1,cost] = gradient_descent(b0,b1,X,y,alpha)
# print('b0 and b1 values are:',b0,b1)

Y = b0+b1*X
plt.plot(X,Y,'r')
plt.show()
# print(cost)
f = np.linspace(0,iterations,5000)
plt.plot(f[:100],cost[:100])

140 140
```



Out[7]: [



In []:

2. (d)

$$\frac{\partial J(b)}{\partial b_0} \Rightarrow -\frac{1}{m} \sum_{i=1}^m (y^{(i)} - b_0 - b_1 x^{(i)}) = 0$$

$$\sum_{i=1}^m (y^{(i)} - b_0 - b_1 x^{(i)}) = 0$$

$$m b_0 = \sum_{i=1}^m y^{(i)} - \sum_{i=1}^m b_1 x^{(i)}$$

$$b_0 = \frac{\sum_{i=1}^m y^{(i)} - \sum_{i=1}^m b_1 x^{(i)}}{m}$$

$$\frac{\partial J(b)}{\partial b_1} = -\frac{1}{m} \sum_{i=1}^m [(y^{(i)} - b_0 - b_1 x^{(i)}) x^{(i)}] = 0$$

$$\Rightarrow \sum_{i=1}^m x^{(i)} y^{(i)} - \sum_{i=1}^m x^{(i)} b_0 - \sum_{i=1}^m b_1 x^{(i)2} = 0$$

$$\Rightarrow \sum_{i=1}^m b_1 x^{(i)2} = \sum_{i=1}^m x^{(i)} y^{(i)} - \sum_{i=1}^m x^{(i)} b_0$$

$$\Rightarrow b_1 = \frac{\sum_{i=1}^m x^{(i)} (y^{(i)} - b_0)}{\sum_{i=1}^m (x^{(i)})^2}$$

Gradient descent is an ~~derivative~~ iterative optimization algorithm used to minimize a function by iteratively moving in the direction of the steepest descent, which is found by taking the negative of the gradient. It is used in machine learning to update parameters of a model, e.g. coefficients in linear regression and weights in neural networks.

The 3 types of gradient descent are: $\sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m} \quad (2) \cdot 5$

- 1) Batch: The weights of the network are updated after the entire dataset (batch) is processed. $\sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$
- 2) Stochastic: The weights of the network are updated for each sample at every step and the weights change very quickly. $\sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$
- 3) Mini-batch: Tradeoff between batch and stochastic gradient descent. Uses a subset of the dataset to update weights at each step. $\sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$

$$0 = \left[\sum_{i=1}^M \frac{1}{m} \right] \sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$$

$$0 = \sum_{i=1}^M \frac{1}{m} \sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$$

$$\sum_{i=1}^M \frac{1}{m} \sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$$

$$\left(\sum_{i=1}^M \frac{1}{m} \right) \sum_{i=1}^M \frac{1}{m} = \frac{(d) \cdot 10}{m}$$

been - through intelligent training in a neural network
 network all in many different ways to achieve a goal
 all parts of the network are connected to each other
 it is possible to have a network that is not fully connected
 in the network, for example, a fully connected network
 network, however, in the network, there is a network.